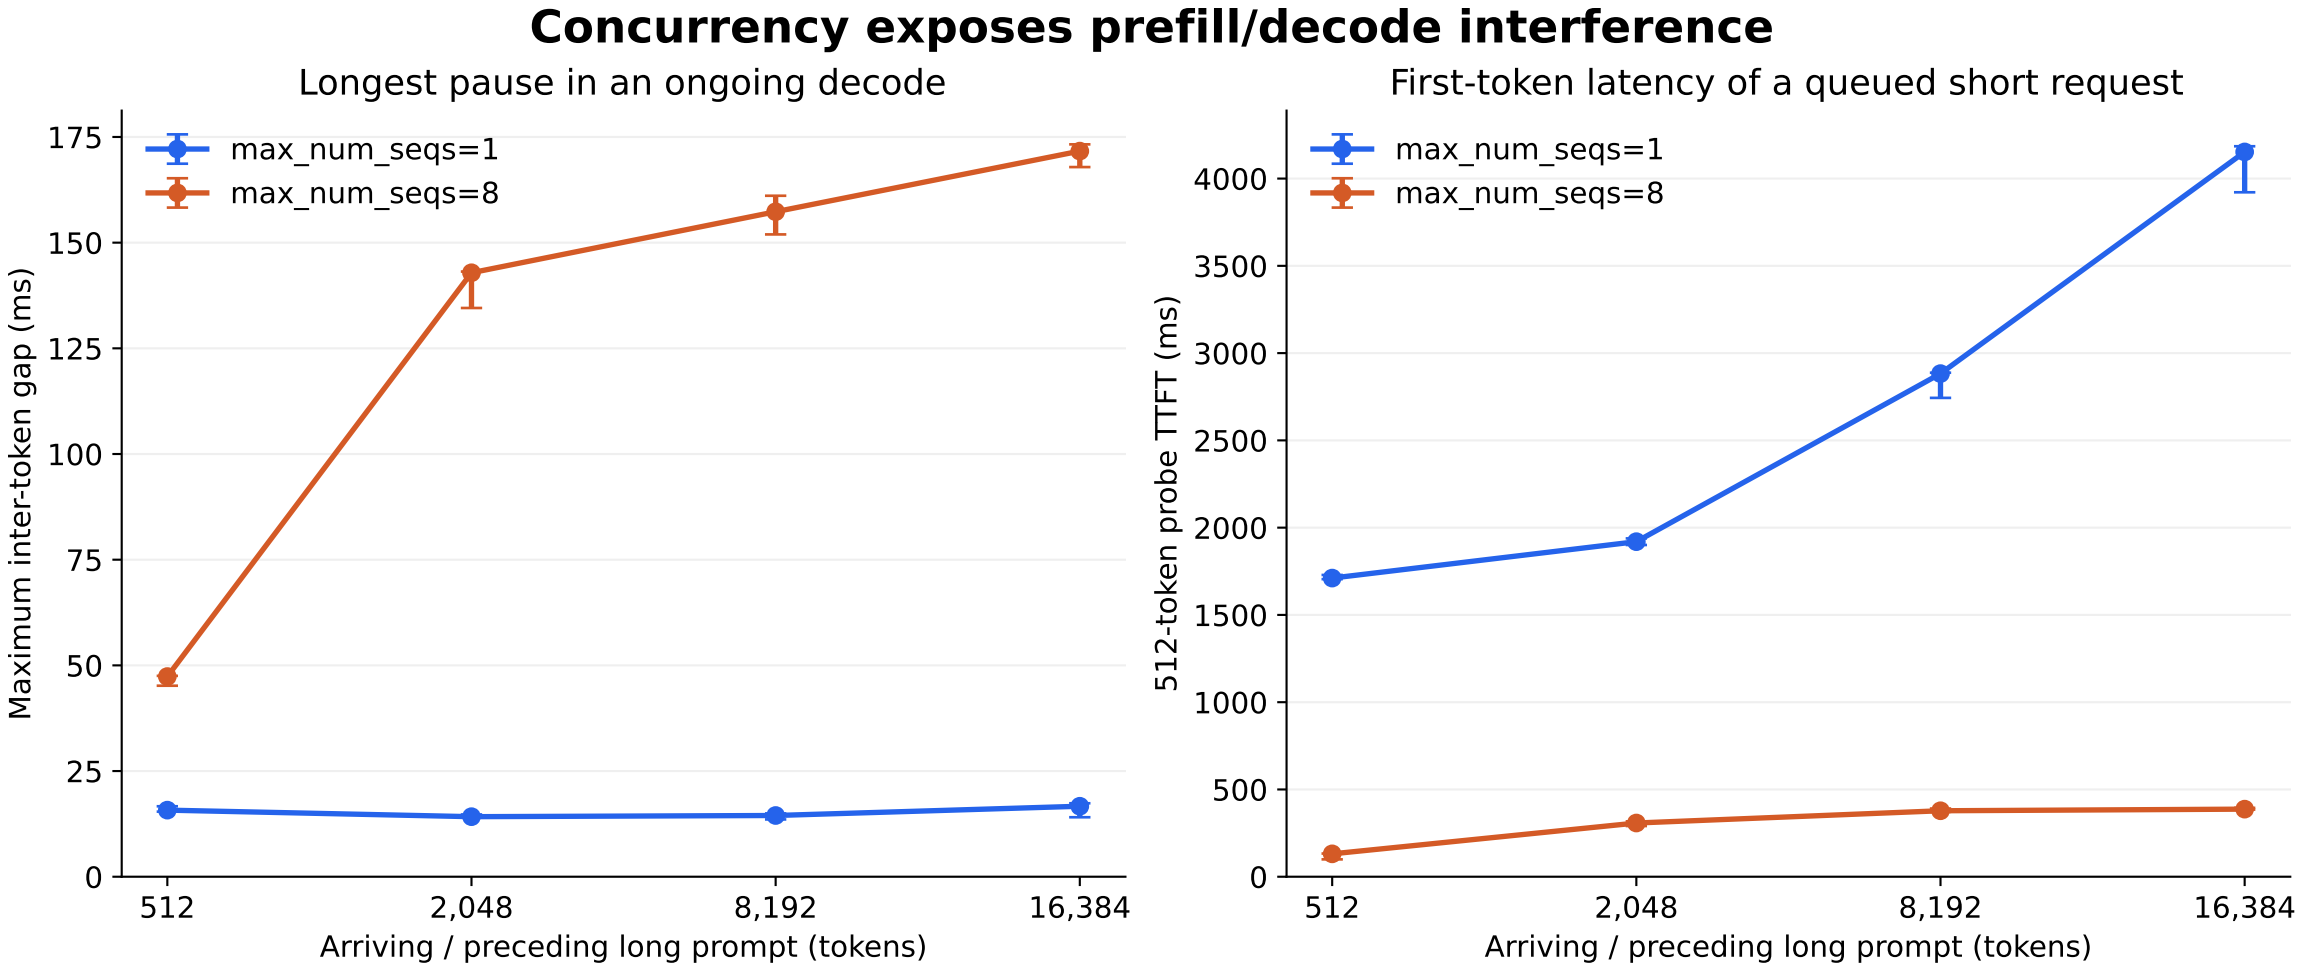
Concurrency exposes prefill/decode interference
Longest pause in an ongoing decode
First-token latency of a queued short request
max_num_seqs=1
max_num_seqs=8
Maximum inter-token gap (ms)
512-token probe TTFT (ms)
0
25
50
75
100
125
150
175
500
1000
1500
2000
2500
3000
3500
4000
512
2,048
8,192
16,384
Arriving / preceding long prompt (tokens)
RTX 6000 Ada · Qwen3.5-9B GPTQ INT4 · FP8 KV · eager · 3 repetitions · median with min–max range
Same model and serving flags; concurrency and port changed. Separate runs; clocks were not locked.